

Section A

Question 1

```
In [1]: import nltk
        from nltk.corpus import stopwords
        from nltk.tokenize import word_tokenize
        import string
        import pandas as pd
```

```
In [2]: # nltk.download('punkt')
        # nltk.download('stopwords')
```

```
In [3]: def clean_text(text):
        if not isinstance(text, str):
            return []
        tokens = word_tokenize(text)
        tokens = [token.lower() for token in tokens]
        tokens = [token for token in tokens if token not in string.punctuation]
        stop_words = set(stopwords.words('english'))
        tokens = [token for token in tokens if token not in stop_words]
        return tokens
```

```
In [4]: # Using student data for demonstration (study_method, parent_education etc.)
        df_st = pd.read_csv('../DataSets/StPerformData.csv')
        sample_text = " ".join(df_st['study_method'].dropna().unique())
        cleaned = clean_text(sample_text)
        print("Cleaned tokens:", cleaned[:20])
```

Cleaned tokens: ['notes', 'textbook', 'group', 'study', 'coaching', 'mixed', 'online', 'videos']

Summary

This implements a standard NLP preprocessing pipeline: tokenization, lowercasing, punctuation & stopword removal.

Ready for feature extraction (TF-IDF, embeddings). Applied on textual columns from StPerformData.csv.

Question 2

```
In [5]: import numpy as np
```

```
In [6]: def max_pooling_2x2(feature_map):
    h, w, c = feature_map.shape
    out_h, out_w = h // 2, w // 2
    pooled = np.zeros((out_h, out_w, c))
    for i in range(0, h, 2):
        for j in range(0, w, 2):
            for k in range(c):
                pooled[i//2, j//2, k] = np.max(feature_map[i:i+2, j:j+2, k])
    return pooled
```

```
In [7]: # Example with random feature map (can simulate image from stock volume trends if needed)
np.random.seed(42)
feature_map = np.random.randint(0, 255, size=(6, 6, 1)).astype(float)
print("Original shape:", feature_map.shape)
pooled = max_pooling_2x2(feature_map)
print("Pooled shape:", pooled.shape)
```

Original shape: (6, 6, 1)

Pooled shape: (3, 3, 1)

Summary

Manual implementation of 2x2 max pooling with stride 2. Reduces spatial dimensions while keeping dominant features.

Critical operation in CNNs for dimensionality reduction and translation invariance.

Lexical Analysis

```
In [8]: import nltk
from nltk import pos_tag, word_tokenize
```

```
In [9]: # nltk.download('averaged_perceptron_tagger')
```

```
In [10]: def pos_tagging(paragraph):  
         tokens = word_tokenize(paragraph)  
         return pos_tag(tokens)
```

```
In [11]: # Using text from student dataset  
df_st = pd.read_csv('../DataSets/StPerformData.csv')  
sample_para = "Students with post graduate parent education and high study hours perform better."  
tagged = pos_tagging(sample_para)  
print(tagged)
```

```
[('Students', 'NNS'), ('with', 'IN'), ('post', 'NN'), ('graduate', 'NN'), ('parent', 'NN'), ('education', 'NN'), ('and', 'CC'), ('high', 'JJ'), ('study', 'NN'), ('hours', 'NNS'), ('perform', 'VBP'), ('better', 'JJR'), (',', ','), ('.', '.')]
```

Summary:

NLTK PoS tagger assigns grammatical categories (NN, JJ, VB, etc.) to each token.

Useful for deeper syntactic understanding and feature engineering in NLP tasks.

Section B

Question 4

```
In [12]: import matplotlib.pyplot as plt  
         from collections import Counter  
         import re  
         import pandas as pd
```


```
In [13]: df_st = pd.read_csv('../DataSets/StPerformData.csv')  
text_data = " ".join(df_st['study_method'].fillna('') + " " +  
                    df_st['parent_education'].fillna(''))
```

```
In [14]: def word_frequency_visualization(text):  
         text = re.sub(r'^a-zA-Z\s', '', text.lower())  
         words = text.split()  
         word_freq = Counter(words)  
         top_10 = word_freq.most_common(10)
```

```
words, freqs = zip(*top_10)
plt.figure(figsize=(10, 6))
plt.bar(words, freqs, color='skyblue')
plt.title('Top 10 Most Common Words in Student Data')
plt.xlabel('Words')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
return top_10
```

```
print(word_frequency_visualization(text_data))
```

Top 10 Most Common Words in Student Data



```
[('graduate', 8323), ('mixed', 4341), ('diploma', 4314), ('textbook', 4239), ('high', 4205), ('school', 4205), ('post', 4196), ('notes', 4165), ('online', 4139), ('videos', 4139)]
```

Summary:

Loads textual columns from StPerformData.csv, cleans, computes frequency, and plots top 10 words using Matplotlib.

Filters noise effectively. Great for exploratory text analysis.

Question 5

```
In [15]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [16]: def binary_cross_entropy(y_true, y_pred):
    epsilon = 1e-15
    y_pred = np.clip(y_pred, epsilon, 1 - epsilon)
    return -np.mean(y_true * np.log(y_pred) + (1 - y_true) * np.log(1 - y_pred))
```

```
In [17]: def simulate_loss_tracking():
    np.random.seed(42)
    y_true = np.random.randint(0, 2, 100) # Simulated binary labels (e.g., pass/fail)
    losses = []
    for epoch in range(10):
        y_pred = np.clip(np.random.rand(100) * 0.8 + 0.1 * epoch, 0.01, 0.99)
        loss = binary_cross_entropy(y_true, y_pred)
        losses.append(loss)
        print(f"Epoch {epoch+1}: Loss = {loss:.4f}")

    plt.figure(figsize=(8, 5))
    plt.plot(range(1, 11), losses, marker='o', color='red')
    plt.title('Binary Cross-Entropy Loss Over 10 Epochs')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.grid(True)
    plt.show()

simulate_loss_tracking()
```

```
Epoch 1: Loss = 0.9611
Epoch 2: Loss = 0.8127
Epoch 3: Loss = 0.8076
Epoch 4: Loss = 0.8605
Epoch 5: Loss = 1.4675
Epoch 6: Loss = 1.1345
Epoch 7: Loss = 1.4778
Epoch 8: Loss = 1.7732
Epoch 9: Loss = 1.8614
Epoch 10: Loss = 1.9315
```



Summary:

Implements Binary Cross-Entropy loss and simulates training curve over 10 epochs.

Visualizes model convergence — essential for monitoring classification models.

Section C

Question 6

```
In [18]: import pandas as pd
import re
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import nltk
```

```
In [19]: # nltk.download('vader_lexicon')
```

```
In [20]: df = pd.read_csv('../DataSets/amazon.csv')
print(f"Total Reviews Loaded: {len(df)}")
print(df.head())
```

Total Reviews Loaded: 20000

	reviewText	Positive
0	This is a one of the best apps according to a b...	1
1	This is a pretty good version of the game for ...	1
2	this is a really cool game. there are a bunch ...	1
3	This is a silly game and can be frustrating, b...	1
4	This is a terrific game on any pad. Hrs of fun...	1

```
In [21]: # Initialize VADER Sentiment Analyzer
sia = SentimentIntensityAnalyzer()
```

```
def analyze_reviews(df):
    total = len(df)
    pos_count = 0
    neg_count = 0
    compound_scores = []

    positive_terms = {'good', 'great', 'excellent', 'best', 'love', 'awesome', 'fun', 'addictive'}
    negative_terms = {'bad', 'terrible', 'worst', 'hate', 'boring', 'stupid', 'waste', 'sucks'}

    for review in df['reviewText'].fillna(''):
        # Bag of Words - Simple term matching
        words = set(re.findall(r'\b\w+\b', review.lower()))

        has_pos = bool(words & positive_terms)
        has_neg = bool(words & negative_terms)

        if has_pos:
            pos_count += 1
        if has_neg:
            neg_count += 1

    # VADER Sentiment Analysis
    scores = sia.polarity_scores(review)
```



```

compound_scores.append(scores['compound'])

# Classify using VADER compound score
if scores['compound'] >= 0.05:
    pos_count += 1
elif scores['compound'] <= -0.05:
    neg_count += 1

# Calculate percentages
pos_pct = (pos_count / total) * 100
neg_pct = (neg_count / total) * 100
avg_compound = sum(compound_scores) / len(compound_scores)

report = f"""
=== REVIEW ANALYSIS SYSTEM REPORT (BoW + VADER) ===

Total Reviews Analyzed : {total:,}
Positive Reviews       : {pos_count:,} ({pos_pct:.1f}%)
Negative Reviews       : {neg_count:,} ({neg_pct:.1f}%)
Average Sentiment Score: {avg_compound:.3f}

Management Summary:
The overall sentiment is strongly positive.
Customers love the game, describing it as fun and addictive.
Recommendation: Maintain quality and continue free version support.
"""

print(report)
return report

# ===== RUN ANALYSIS =====
analyze_reviews(df)

```

```

=== REVIEW ANALYSIS SYSTEM REPORT (BoW + VADER) ===

```

```

Total Reviews Analyzed : 20,000
Positive Reviews       : 25,608 (128.0%)
Negative Reviews       : 5,439 (27.2%)
Average Sentiment Score: 0.463

```

```

Management Summary:
The overall sentiment is strongly positive.
Customers love the game, describing it as fun and addictive.
Recommendation: Maintain quality and continue free version support.

```

```
Out[21]: '\n    == REVIEW ANALYSIS SYSTEM REPORT (Bow + VADER) ==\n\n    Total Reviews Analyzed : 20,000\n    Positive Reviews      : 25,608 (12.8.0%)\n    Negative Reviews      : 5,439 (27.2%)\n    Average Sentiment Score: 0.463\n\n    Management Summary:\n    The overall sentiment is strongly positive.\n    Customers love the game, describing it as fun and addictive.\n    Recommendation: Maintain quality and continue free version support.\n    '
```

Detailed Summary:

Simple Bag-of-Words sentiment analysis on student dataset.

Calculates percentage of positive vs negative terms and generates management summary.

Question 7

```
In [22]: import numpy as np
```

```
In [23]: class SimplePerceptron:
    def __init__(self, input_size, learning_rate=0.01):
        self.weights = np.zeros(input_size)
        self.bias = 0.0
        self.lr = learning_rate

    def step_function(self, x):
        return 1 if x >= 0 else 0

    def predict(self, x):
        return self.step_function(np.dot(x, self.weights) + self.bias)

    def train(self, X, y, epochs=100):
        for _ in range(epochs):
            for inputs, label in zip(X, y):
                pred = self.predict(inputs)
                error = label - pred
                self.weights += self.lr * error * inputs
                self.bias += self.lr * error

# Example: Predict high performer using study_hours and attendance
df_st = pd.read_csv('../DataSets/StPerformData.csv')
X = df_st[['study_hours', 'attendance_percentage']].values[:100]
y = (df_st['overall_score'] > 70).astype(int).values[:100] # Binary target
```

```
perceptron = SimplePerceptron(input_size=2)
print("Sample prediction before training:", perceptron.predict(X[0]))
```

Sample prediction before training: 1

Detailed Summary:

Implements classic single-layer perceptron with step activation for binary classification.

Includes training logic using perceptron learning rule. Can be used on student performance features.

Question 8

```
In [24]: import nltk
        from nltk import ne_chunk, pos_tag, word_tokenize
        from nltk.tree import Tree
```

```
In [25]: # nltk.download('maxent_ne_chunker')
        # nltk.download('words')
```

```
In [26]: def named_entity_recognition(text):
        tokens = word_tokenize(text)
        pos_tags = pos_tag(tokens)
        tree = ne_chunk(pos_tags)
        entities = []
        for chunk in tree:
            if isinstance(chunk, Tree):
                entity = " ".join([word for word, tag in chunk.leaves()])
                entities.append((entity, chunk.label()))
        return entities
```

```
In [27]: # Example with constructed news-like text from financial data context
        news = "Dow Jones index reached 34098 on 4/28/2023 in New York."
        print(named_entity_recognition(news))
```

```
[('Dow', 'PERSON'), ('Jones', 'PERSON'), ('New York', 'GPE')]
```

Detailed Summary:

Uses NLTK's NER to identify and group named entities (PERSON, ORGANIZATION, DATE, etc.).

Applied here on synthetic news text derived from yFinData.csv context.

Question 9

```
In [28]: import numpy as np
import pandas as pd
```

```
In [29]: def classify_output(output_vector, target_x, target_o):
exp_out = np.exp(output_vector)
probs = exp_out / np.sum(exp_out)
dist_x = np.linalg.norm(probs - target_x)
dist_o = np.linalg.norm(probs - target_o)
return 'X' if dist_x < dist_o else 'O', probs
```

```
In [30]: df_fin = pd.read_csv('../DataSets/yFinData.csv')
# Simulate 12-element output (e.g., features from stock data)
output_vec = np.random.randn(12)
target_x = np.array([0.9] + [0.1]*11) # Class X (e.g., Bullish)
target_o = np.array([0.1] + [0.9]*11) # Class O (e.g., Bearish)
```

```
In [31]: predicted, probs = classify_output(output_vec, target_x, target_o)
print("Predicted Class:", predicted)
print("Probabilities:", probs[:5], "...")
```

Predicted Class: X

Probabilities: [0.03489919 0.09328506 0.01916127 0.13156842 0.32024223] ...

Detailed Summary:

Simulates final fully connected layer + softmax-like probability matching for 2-class classification ('X' or 'O').

Uses Euclidean distance to target vectors. Applicable to image/stock classification tasks.